

MACHINE LEARNING

LAB EXPERIEMENTS

SUFYAAN EXP 1

```
# a) Operation of data types in Python.

# Create a variable of type integer.
print("experiement 1a")
x = 10

# Create a variable of type float.
y = 10.5

# Create a variable of type string.
z = "Hello, world!"

# Print the values of the variables.
print(x)
print(y)
print(z)

# b) Different Arithmetic Operations on numbers in Python.
print("\n experiement 1b")
# Add two numbers.
x = 10
y = 5

z = x + y

# Print the sum of the two numbers.
print(z)

# Subtraction
d = x - y
print("Subtraction:", d)

# Multiplication
e = x * y
print("Multiplication:", e)

# Division
f = x / y
print("Division:", f)

# Modulo
g = x % y
print("Modulo:", g)
```

```
# Exponentiation
h = x ** y
print("Exponentiation:", h)

# Floor division
i = x // y
print("Floor division:", i)
# c) Create, concatenate and print a string and access substring from
a given string.
print("\n experiement 1c")
# Create a string.
str = "Hello, world!"

# Print the string.
print(str)

# Concatenate two strings.
str1 = "Hello"
str2 = "world!"

str = str1 + str2

# Print the concatenated string.
print(str)

# Access a substring from a string.
str = "Hello, world!"

substring = str[0:5]

# Print the substring.
print(substring)

# d) Append, and remove lists in python.
print("\n experiement 1d")
# Create a list.
list = [1, 2, 3, 4, 5]

# Append an element to the list.
list.append(6)

# Print the list.
print(list)

# Remove an element from the list.
list.remove(3)

# Print the list.
print(list)

# e) Demonstrate working with tuples in python.
```

```

print("\n experiement 1e")
# Create a tuple.
tuple = (1, 2, 3, 4, 5)

# Print the tuple.
print(tuple)

# Access an element from a tuple.
element = tuple[0]

# Print the element.
print(element)

# f) Demonstrate working with dictionaries in python.
print("\n experiement 1f")
# Create a dictionary.
dict = {"name": "John", "age": 30}

# Print the dictionary.
print(dict)

# Access a value from a dictionary.
value = dict["name"]

# Print the value.
print(value)

dict["occupation"] = 'engineer'
del dict['name']
print(dict)

```

```

experiement 1a
10
10.5
Hello, world!

```

```

    experiement 1b
15
Subtraction: 5
Multiplication: 50
Division: 2.0
Modulo: 0
Exponentiation: 100000
Floor division: 2

```

```

    experiement 1c
Hello, world!
Helloworld!
Hello

```

```

    experiement 1d

```

```

[1, 2, 3, 4, 5, 6]
[1, 2, 4, 5, 6]

    experiement 1e
(1, 2, 3, 4, 5)
1

    experiement 1f
{'name': 'John', 'age': 30}
John
{'age': 30, 'occupation': 'engineer'}

```

SUFYAAN EXP2

```

import numpy as np

# Define the two arrays.
array1 = np.array([1, 2, 3, 4, 5])
array2 = np.array([6, 7, 8, 9, 10])

# Compute the expected value.
expected_value = np.mean(array1)

# Compute the mean.
mean = np.mean(array1)

# Compute the standard deviation.
standard_deviation = np.std(array1)

# Compute the variance.
variance = np.var(array1)

# Compute the covariance.
covariance = np.cov(array1, array2)[0][1]

# Compute the covariance matrix.
covariance_matrix = np.cov(array1, array2)

# Print the results.
print("Expected value:", expected_value)
print("Mean:", mean)
print("Standard deviation:", standard_deviation)
print("Variance:", variance)
print("Covariance:", covariance)
print("Covariance matrix:", covariance_matrix)

Expected value: 3.0
Mean: 3.0
Standard deviation: 1.4142135623730951
Variance: 2.0

```

Covariance: 2.5
Covariance matrix: $\begin{bmatrix} 2.5 & 2.5 \\ 2.5 & 2.5 \end{bmatrix}$

SUFYAAN EXP3

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn

# Path to the CSV file
csv_file = 'data.csv'

# Read the CSV file into a DataFrame
df = pd.read_csv(csv_file)

# Print the contents of the DataFrame
print(df)

def find_missing_numbers(numbers):
    missing_numbers = []
    for i in range(1, len(numbers) + 1):
        if i not in numbers:
            missing_numbers.append(i)
    return missing_numbers

# Example usage
data = [1, 3, 5, 6, 8, 10]
missing = find_missing_numbers(data)
print("Missing numbers:", missing)
from sklearn.preprocessing import LabelEncoder

# Categorical data
data = ['red', 'blue', 'green', 'blue', 'red', 'green', 'green']

# Create an instance of LabelEncoder
label_encoder = LabelEncoder()

# Fit and transform the data
encoded_data = label_encoder.fit_transform(data)

# Print the encoded data
print("Encoded Data:", encoded_data)

# Print the original classes
print("Original Classes:", label_encoder.classes_)
import random
```

```

# Original dataset
dataset = ['data_point1', 'data_point2', 'data_point3', 'data_point4',
'data_point5']

# Set the seed for reproducibility (optional)
random.seed(42)

# Specify the split ratio (e.g., 80% training, 20% testing)
split_ratio = 0.8

# Calculate the number of training samples based on the split ratio
num_train_samples = int(len(dataset) * split_ratio)

# Shuffle the dataset randomly
random.shuffle(dataset)

# Split the dataset into training and test data
train_data = dataset[:num_train_samples]
test_data = dataset[num_train_samples:]

# Print the training and test data
print("Training Data:")
print(train_data)
print("\nTest Data:")
print(test_data)
import numpy as np

# Input data
data = np.array([[2, 1000], [5, 2000], [10, 5000], [3, 1500]])

# Min-Max Scaling
min_val = np.min(data, axis=0)
max_val = np.max(data, axis=0)
scaled_data = (data - min_val) / (max_val - min_val)

print("Min-Max Scaling:")
print(scaled_data)

# Standardization
mean_val = np.mean(data, axis=0)
std_val = np.std(data, axis=0)
standardized_data = (data - mean_val) / std_val

print("\nStandardization:")
print(standardized_data)

```

	23	0	63	90	20	100	30.4	0.25	1
0	20	10	110	80	40	140	26.4	0.20	1
1	25	5	140	70	30	120	33.6	0.28	0
2	30	8	110	60	20	90	25.2	0.35	0

```

3    35    12   120    70    40   110   30.8   0.40   0
4    40     4    90    50    20    80   23.4   0.38   0
5    45     2   120    60    30   100   28.0   0.45   0
6    50     6   130    70    40   130   32.4   0.48   0
7    55     9    80    50    20    70   21.6   0.32   0
8    60    11   100    60    30    90   26.4   0.38   0
9    65    15   160    90    30   180   38.0   0.59   1
10   70    17   190   100    35   200   42.4   0.68   1
11   75    19   220   110    40   220   46.8   0.77   1
12   80    21   250   120    45   240   51.2   0.86   1
13   85    23   280   130    50   260   55.6   0.95   1
14   90    25   310   140    55   280   60.0   1.04   1
15   95    27   340   150    60   300   64.4   1.13   1
16  100    29   370   160    65   320   68.8   1.22   1
17  105    31   400   170    70   340   73.2   1.31   1
18  110    33   430   180    75   360   77.6   1.40   1
Missing numbers: [2, 4]
Encoded Data: [2 0 1 0 2 1 1]
Original Classes: ['blue' 'green' 'red']
Training Data:
['data_point4', 'data_point2', 'data_point3', 'data_point5']

Test Data:
['data_point1']
Min-Max Scaling:
[[0.    0.   ]
 [0.375 0.25 ]
 [1.    1.   ]
 [0.125 0.125]]

Standardization:
[[-0.97332853 -0.88354126]
 [ 0.         -0.2409658 ]
 [ 1.62221421  1.68676059]
 [-0.64888568 -0.56225353]]

```

SUFYAAN EXP 4

```

import pandas as pd
import numpy as np

# Function to generate synthetic diabetes dataset with 100 rows
def generate_synthetic_dataset(num_rows=100):
    np.random.seed(42)

    # Generate random values for each attribute based on some ranges
    feature1 = np.random.randint(20, 70, num_rows)
    feature2 = np.random.randint(0, 15, num_rows)
    feature3 = np.random.randint(60, 200, num_rows)

```

```

target_variable = np.random.randint(50, 100, num_rows)

# Create the synthetic dataset
dataset = pd.DataFrame({
    'feature1': feature1,
    'feature2': feature2,
    'feature3': feature3,
    'target_variable': target_variable
})

return dataset

# Generate the synthetic dataset with 100 rows
dataset = generate_synthetic_dataset(num_rows=100)

# Save the dataset to a CSV file
dataset.to_csv('dataset.csv', index=False)

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt

# Load your dataset into a pandas DataFrame
# Assuming you have a CSV file named 'data.csv' with features X and
# target variable y
data = pd.read_csv('dataset.csv')

# Split the data into features (X) and target variable (y)
X = data[['feature1', 'feature2', 'feature3']] # Replace 'feature1',
'feature2', 'feature3' with your actual feature names
y = data['target_variable'] # Replace 'target_variable' with your
actual target variable name

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42) # Adjust test_size as desired

# Create a linear regression model
model = LinearRegression()
# Train the model using the training data
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)

```



```

# Compare actual values with predicted values
result_df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
print(result_df)

# Plot the training set
plt.scatter(X_train['feature1'], y_train, color='blue',
            label='Training Set')
# Plot the test set
plt.scatter(X_test['feature1'], y_test, color='red', label='Test Set')
# Plot the best-fit line
# plt.plot(X_train['feature1'], model.predict(X_train), color='black',
#          label='Best Fit')

plt.xlabel('Feature 1')
plt.ylabel('Target Variable')
plt.title('Training Set vs. Test Set')
plt.legend()
plt.show()

# Make predictions on the test set
y_pred = model.predict(X_test)


# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Print the evaluation metrics
print('Mean Squared Error:', mse)
print('R-squared:', r2)

```

	Actual	Predicted
83	77	78.569972
53	91	76.837071
70	62	73.285512
45	73	81.989229
44	95	75.673908
39	74	75.428690
22	55	72.829900
80	61	73.984318
10	54	76.826774
0	57	76.984376
18	77	74.738071
30	80	75.733680
73	94	76.136241
33	70	76.251352
90	71	74.311676
4	70	73.686093
76	57	75.263717

77	95	77.680035
12	68	79.384234
31	84	78.143493



Mean Squared Error: 169.86028188186393
R-squared: 0.013515604315853857

SUFYAAN EXP 5

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# Load the dataset
data = pd.read_excel('iml exp 5.xlsx') # Replace 'your_dataset.csv'
with the actual filename

# Display the first few rows of the dataset
print(data.head())
# Split the dataset into features (X) and target variable (y)
```

```

X = data[['Gender', 'Age', 'EstimatedSalary']] # Replace with
relevant feature columns
y = data['Purchased'] # Replace with the target variable column

# Convert categorical variable 'Gender' into numeric using one-hot
encoding
X = pd.get_dummies(X, drop_first=True)

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
# Split the dataset into features (X) and target variable (y)
X = data[['Gender', 'Age', 'EstimatedSalary']] # Replace with
relevant feature columns
y = data['Purchased'] # Replace with the target variable column

# Convert categorical variable 'Gender' into numeric using one-hot
encoding
X = pd.get_dummies(X, drop_first=True)

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
# Create a logistic regression model
model = LogisticRegression()

# Train the model using the training data
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print('Accuracy:', accuracy)

```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

Accuracy: 0.75

SUFYAAN EXP 6A

```

probAbsentFriday=0.03
probFriday=0.2

```

```

# bayes Formula
#p(Absent|Friday)=p(Friday|Absent)p(Absent)/p(Friday)
#p(Friday|Absent)=p(FridaynAbsent)/p(Absent)
# Therefore the result is:
bayesResult=(probAbsentFriday/probFriday)
print(bayesResult * 100)

15.0

```

SUFYAAN EXP 6B

```

import csv
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score
# Step 1: Load the data
def load_data(filename):
    with open(filename, 'r') as csvfile:
        data = list(csv.reader(csvfile))
    return data
# Step 2: Preprocess the data (if needed)
# Step 3: Split the data
import numpy as np
def split_data(data, test_size):
    X = np.array([row[:-1] for row in data], dtype=np.float64)
    y = [row[-1] for row in data]
    X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=test_size, random_state=42)
    return X_train, X_test, y_train, y_test
# Step 4: Train the Naïve Bayesian classifier
def train_naive_bayes(X_train, y_train):
    clf = GaussianNB()
    clf.fit(X_train, y_train)
    return clf
# Step 5: Make predictions
def predict(clf, X_test):
    return clf.predict(X_test)
# Step 6: Evaluate accuracy
def evaluate_accuracy(y_true, y_pred):
    return accuracy_score(y_true, y_pred)
# Main function
if __name__ == "__main__":
    filename = "data.csv"
    data = load_data(filename)
    # Set the test size for splitting the data
    test_size = 0.2
    X_train, X_test, y_train, y_test = split_data(data, test_size)
    clf = train_naive_bayes(X_train, y_train)

```

```

y_pred = predict(clf, X_test)
accuracy = evaluate_accuracy(y_test, y_pred)
print(f"Accuracy: {accuracy}")

```

Accuracy: 0.5

SUFYAAN EXP 7

```

import pandas as pd
import numpy as np
import csv
from pgmpy.models import BayesianNetwork
from pgmpy.estimators import MaximumLikelihoodEstimator
from pgmpy.inference import VariableElimination

# Step 1: Load and preprocess the dataset
data = pd.read_csv("archive/diabetes.csv")
# data = data.replace('?', np.nan)

# data = data.apply(pd.to_numeric)
print('Few examples from the dataset are given below')
print(data.head())

# Step 2: Define the structure of the Bayesian network
from pgmpy.estimators import MaximumLikelihoodEstimator

model = BayesianNetwork([
    ('Age', 'Outcome'),
    ('BMI', 'Outcome'),
    ('BloodPressure', 'Outcome'),
    ('Insulin', 'Outcome'),
    ('Glucose', 'Outcome'),
    ('Pregnancies', 'Outcome')
])

# Step 3: Learn the parameters of the Bayesian network
print('\n Learning CPD using Maximum likelihood estimators')
# model.fit(data)
model.fit(data, estimator=MaximumLikelihoodEstimator)

# estimator = MaximumLikelihoodEstimator(model, data)
# model.fit(data, estimator)

# Step 4: Perform diabetic patient detection
print('\n Inferencing with Bayesian Network:')

print('\n Inferencing with Bayesian Network:')

inference = VariableElimination(model)
# q=inference.query(variables=["Outcome"],evidence={'Age': 34})

```

```

# Sample test data for inference (you can use real test data here)
test_data = {
    'Age': 50,
    'BMI': 31,
    'BloodPressure': 1,
    'Insulin': 3,
    'Glucose': 148,
    'Pregnancies': 1
}

# Query to detect Diabetes (set 'Outcome' to None to calculate its value)
result = inference.map_query(variables=['Outcome'],
                             evidence=test_data)

# Print the result
# print("Result: Diabetic" if result['Outcome'] == 1 else "Result: Non-diabetic")
print("Result: Diabetic" if result['Outcome'] else "Result: Non-diabetic")
# Print the result
print("Result: Diabetic" if result['Outcome'] == 1 else "Result: Non-diabetic")

```

Few examples from the dataset are given below

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	
BMI \						
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

Learning CPD using Maximum likelihood estimators

Inferencing with Bayesian Network:

Inferencing with Bayesian Network:

```
c:\Python310\lib\site-packages\pgmpy\factors\discrete\
DiscreteFactor.py:541: UserWarning: Found unknown state name. Trying
to switch to using all state names as state numbers
```

```
warnings.warn(
Finding Elimination Order: : : 0it [00:00, ?it/s]
0it [00:00, ?it/s]
```

Result: Non-diabetic

Result: Non-diabetic

#Step 1: Import necessary libraries

```
import pandas as pd
from pgmpy.models import BayesianModel
from pgmpy.estimators import MaximumLikelihoodEstimator
from pgmpy.inference import VariableElimination
```

#Step 2: Load and preprocess the dataset

```
data = pd.read_csv("diabetes_data.csv")
```

#Step 3: Define the structure of the Bayesian network

Define the structure of the Bayesian network

```
model = BayesianModel([
    ('Age', 'Diabetes'),
    ('BMI', 'Diabetes'),
    ('BloodPressure', 'Diabetes'),
    ('Insulin', 'Diabetes'),
    ('Glucose', 'Diabetes'),
    ('Pregnancies', 'Diabetes'),
    ('DiabetesPedigreeFunction', 'Diabetes')
])
```

#Step 4: Learn the parameters of the Bayesian network

Learn the parameters of the Bayesian network using Maximum Likelihood Estimation

```
estimator = MaximumLikelihoodEstimator(model, data)
model.fit(data, estimator)
```

#Step 5: Perform diabetic patient detection

Perform inference using Variable Elimination

```
inference = VariableElimination(model)
```

Sample test data for inference (you can use real test data here)

```
test_data = {
    'Age': 40,
    'BMI': 25,
    'BloodPressure': 70,
    'Insulin': 100,
    'Glucose': 140,
    'Pregnancies': 0,
    'DiabetesPedigreeFunction': 0.2
}
```

Query to detect Diabetes (set 'Diabetes' to None to calculate its value)

```
result = inference.map_query(variables=['Diabetes'],
evidence=test_data)
```

```
# Print the result
print("Result: Diabetic" if result['Diabetes'] else "Result: Non-
diabetic")
```

```
c:\Python310\lib\site-packages\tqdm\auto.py:22: TqdmWarning: IProgress
not found. Please update jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
```

```
from .autonotebook import tqdm as notebook_tqdm
c:\Python310\lib\site-packages\pgmpy\models\BayesianModel.py:8:
FutureWarning: BayesianModel has been renamed to BayesianNetwork.
Please use BayesianNetwork class, BayesianModel will be removed in
future.
```

```
warnings.warn(
```

```
-----
-----
TypeError                                Traceback (most recent call
last)
```

```
f:\sufyaaneng_sem4\ml\New folder\labprog\lab.ipynb Cell 6 in 2
    <a href='vscode-notebook-cell:/f%3A/sufyaaneng_sem4/ml/New
%20folder/labprog/lab.ipynb#W5sZmlsZQ%3D%3D?line=18'>19</a> #Step 4:
Learn the parameters of the Bayesian network
```

```
    <a href='vscode-notebook-cell:/f%3A/sufyaaneng_sem4/ml/New
%20folder/labprog/lab.ipynb#W5sZmlsZQ%3D%3D?line=19'>20</a> # Learn
the parameters of the Bayesian network using Maximum Likelihood
Estimation
```

```
    <a href='vscode-notebook-cell:/f%3A/sufyaaneng_sem4/ml/New
%20folder/labprog/lab.ipynb#W5sZmlsZQ%3D%3D?line=20'>21</a> estimator
= MaximumLikelihoodEstimator(model, data)
```

```
--> <a href='vscode-notebook-cell:/f%3A/sufyaaneng_sem4/ml/New
%20folder/labprog/lab.ipynb#W5sZmlsZQ%3D%3D?line=21'>22</a>
model.fit(data, estimator)
```

```
    <a href='vscode-notebook-cell:/f%3A/sufyaaneng_sem4/ml/New
%20folder/labprog/lab.ipynb#W5sZmlsZQ%3D%3D?line=22'>23</a> #Step 5:
Perform diabetic patient detection
```

```
    <a href='vscode-notebook-cell:/f%3A/sufyaaneng_sem4/ml/New
%20folder/labprog/lab.ipynb#W5sZmlsZQ%3D%3D?line=23'>24</a> # Perform
inference using Variable Elimination
```

```
    <a href='vscode-notebook-cell:/f%3A/sufyaaneng_sem4/ml/New
%20folder/labprog/lab.ipynb#W5sZmlsZQ%3D%3D?line=24'>25</a> inference
= VariableElimination(model)
```

```
File c:\Python310\lib\site-packages\pgmpy\models\
BayesianNetwork.py:577, in BayesianNetwork.fit(self, data, estimator,
state_names, complete_samples_only, n_jobs, **kwargs)
```

```
    575     estimator = MaximumLikelihoodEstimator
```

```
    576 else:
```

```
--> 577     if not issubclass(estimator, BaseEstimator):
```

```
    578         raise TypeError("Estimator object should be a valid
pgmpy estimator.")
```



```

580 _estimator = estimator(
581     self,
582     data,
583     state_names=state_names,
584     complete_samples_only=complete_samples_only,
585 )

```

TypeError: issubclass() arg 1 must be a class

SUFYAAN EXP 8

```

import pandas as pd
import numpy as np
# Sample dataset
data = {
    'Can Fly': ['Yes', 'Yes', 'No', 'No', 'Yes', 'No'],
    'Has Fur': ['No', 'Yes', 'Yes', 'Yes', 'No', 'Yes'],
    'Is Mammal': ['No', 'Yes', 'Yes', 'Yes', 'Yes', 'No']
}
df = pd.DataFrame(data)
def entropy(target_col):
    elements, counts = np.unique(target_col, return_counts=True)
    probabilities = counts / counts.sum()
    entropy_val = -sum(probabilities * np.log2(probabilities))
    return entropy_val
def information_gain(data, split_feature, target_feature):
    total_entropy = entropy(data[target_feature])
    # Calculate the weighted entropy for the two splits
    unique_values = data[split_feature].unique()
    weighted_entropy = 0
    for value in unique_values:
        subset = data[data[split_feature] == value]
        subset_entropy = entropy(subset[target_feature])
        weighted_entropy += (len(subset) / len(data)) * subset_entropy
    information_gain_val = total_entropy - weighted_entropy
    return information_gain_val
def build_tree(data, original_data, features, target_feature,
parent_node_class=None):
    # Base cases
    if len(np.unique(data[target_feature])) == 1:
        return np.unique(data[target_feature])[0]

    elif len(data) == 0:
        return np.unique(original_data[target_feature])
    [np.argmax(np.unique(original_data[target_feature],
        return_counts=True)[1])]
    elif len(features) == 0:
        return parent_node_class

```

```

else:
    parent_node_class = np.unique(data[target_feature])
    [np.argmax(np.unique(data[target_feature],
return_counts=True)[1])]
    # Select the feature with the highest information gain
    item_values = [information_gain(data, feature, target_feature) for
feature in features]
    best_feature_index = np.argmax(item_values)
    best_feature = features[best_feature_index]
    # Create the tree structure
    tree = {best_feature: {}}
    # Remove the selected feature from the features list
    features = [f for f in features if f != best_feature]
    # Recurse on the subsets
    for value in np.unique(data[best_feature]):
        subset = data[data[best_feature] == value]
        subtree = build_tree(subset, data, features, target_feature,
parent_node_class)
        tree[best_feature][value] = subtree
    return tree
# Building the decision tree
features = ['Can Fly', 'Has Fur']
target_feature = 'Is Mammal'
tree = build_tree(df, df, features, target_feature)
print(tree)

{'Has Fur': {'No': {'Can Fly': {'Yes': 'No'}}}}

```

SUFYAAN EXP 9

```

import numpy as np
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
# Load the Iris dataset
iris = load_iris()
X, y = iris.data, iris.target
# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
# Standardize the features (optional but recommended for KNN)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
# Define the value of k for KNN
k = 3

```

```

# Create the KNN classifier and fit it to the training data
knn_classifier = KNeighborsClassifier(n_neighbors=k)
knn_classifier.fit(X_train, y_train)
# Make predictions on the test set
y_pred = knn_classifier.predict(X_test)
# Calculate and print the accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")

```

Accuracy: 1.00

SUFYAAN EXP 10

```

from sklearn.cluster import KMeans
import numpy as np
X = np.array([[1.713,1.586], [0.180,1.786], [0.353,1.240],
[0.940,1.566], [1.486,0.759],
[1.266,1.106],[1.540,0.419],[0.459,1.799],[0.773,0.186]])
y=np.array([0,1,1,0,1,0,1,1,1])
kmeans = KMeans(n_clusters=3, random_state=0).fit(X,y)
print("The input data is ")
print("VAR1 \t VAR2 \t CLASS")
i=0
for val in X:
    print(val[0],"\t",val[1],"\t",y[i])
    i+=1
    print("="*20)
# To get test data from the user
print("The Test data to predict ")
test_data = []
VAR1 = float(input("Enter Value for VAR1 :"))
VAR2 = float(input("Enter Value for VAR2 :"))
test_data.append(VAR1)
test_data.append(VAR2)
print("="*20)
print("The predicted Class is : ", kmeans.predict([test_data]))

c:\Python310\lib\site-packages\sklearn\cluster\_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to
'auto' in 1.4. Set the value of `n_init` explicitly to suppress the
warning
    warnings.warn(

```

```

The input data is
VAR1  VAR2  CLASS
1.713    1.586    0
=====
0.18  1.786    1
=====
0.353    1.24    1

```

```

=====
0.94   1.566       0
=====
1.486       0.759       1
=====
1.266       1.106       0
=====
1.54   0.419       1
=====
0.459       1.799       1
=====
0.773       0.186       1
=====
The Test data to predict
=====
The predicted Class is :  [1]

```

SUFYAAN EXP 11

```

# Sample data (features: study hours, previous test scores; label:
pass/fail)
import numpy as np
np.random.seed(42)
study_hours = np.random.randint(1, 8, 100)
test_scores = np.random.randint(40, 100, 100)

# Generate the target variable 'pass_fail' with a mix of class labels
pass_fail = np.zeros(100)
pass_indices = np.random.choice(100, 30, replace=False)
pass_fail[pass_indices] = 1

# Split data into training and testing sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    np.column_stack((study_hours, test_scores)),
    pass_fail, test_size=0.2, random_state=42)

# Train different classifiers and evaluate their performance
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
# Logistic Regression
logreg = LogisticRegression()
logreg.fit(X_train, y_train)
y_pred_logreg = logreg.predict(X_test)
# Decision Tree
dtree = DecisionTreeClassifier()
dtree.fit(X_train, y_train)

```

```

y_pred_dtree = dtree.predict(X_test)
# Random Forest
rf = RandomForestClassifier()
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
# Evaluate metrics
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, roc_auc_score
# Logistic Regression
print("Logistic Regression:")
print("Accuracy:", accuracy_score(y_test, y_pred_logreg))
print("Precision:", precision_score(y_test, y_pred_logreg))
print("Recall:", recall_score(y_test, y_pred_logreg))
print("F1-score:", f1_score(y_test, y_pred_logreg))
print("AUC-ROC:", roc_auc_score(y_test, y_pred_logreg))
print()
# Decision Tree
print("Decision Tree:")
print("Accuracy:", accuracy_score(y_test, y_pred_dtree))
print("Precision:", precision_score(y_test, y_pred_dtree))
print("Recall:", recall_score(y_test, y_pred_dtree))
print("F1-score:", f1_score(y_test, y_pred_dtree))
print("AUC-ROC:", roc_auc_score(y_test, y_pred_dtree))
print()
# Random Forest
print("Random Forest:")
print("Accuracy:", accuracy_score(y_test, y_pred_rf))
print("Precision:", precision_score(y_test, y_pred_rf))
print("Recall:", recall_score(y_test, y_pred_rf))
print("F1-score:", f1_score(y_test, y_pred_rf))
print("AUC-ROC:", roc_auc_score(y_test, y_pred_rf))

```

```

Logistic Regression:
Accuracy: 0.65
Precision: 0.0
Recall: 0.0
F1-score: 0.0
AUC-ROC: 0.4642857142857143

```

```

Decision Tree:
Accuracy: 0.55
Precision: 0.2
Recall: 0.16666666666666666
F1-score: 0.1818181818181818
AUC-ROC: 0.44047619047619047

```

```

Random Forest:
Accuracy: 0.5
Precision: 0.16666666666666666
Recall: 0.16666666666666666

```

F1-score: 0.16666666666666666
AUC-ROC: 0.40476190476190477

SUFYAAN EXP 12

```
# Importing necessary libraries
import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix
# Load the Iris dataset and consider only two classes: setosa and
versicolor
iris = datasets.load_iris()
X = iris.data
y = iris.target
X = X[y != 2, :2] # Consider only setosa and versicolor classes
y = y[y != 2]
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)
# Create and train the SVM classifier
svm_classifier = SVC(kernel='linear', C=1.0)
svm_classifier.fit(X_train, y_train)
# Predict on the test set
y_pred = svm_classifier.predict(X_test)
# Calculate the accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
# Create a confusion matrix to evaluate the performance
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(conf_matrix)
# Plot the decision boundary
# Function to plot the decision boundary
def plot_decision_boundary(clf, X, y, title):
    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.coolwarm, s=30)
    ax = plt.gca()
    xlim = ax.get_xlim()
    ylim = ax.get_ylim()

    # Create grid to evaluate the model
    xx = np.linspace(xlim[0], xlim[1], 30)
    yy = np.linspace(ylim[0], ylim[1], 30)
    YY, XX = np.meshgrid(yy, xx)
    xy = np.vstack([XX.ravel(), YY.ravel()]).T
    Z = clf.decision_function(xy).reshape(XX.shape)
```

```

# Plot decision boundary and margins
ax.contour(XX, YY, Z, colors='k', levels=[-1, 0, 1], alpha=0.5,
linestyles=['--', '-', '--'])


# Plot support vectors
ax.scatter(clf.support_vectors[:, 0], clf.support_vectors[:, 1],
s=100,
linewidth=1, facecolors='none', edgecolors='k')

plt.xlabel('Sepal Length (cm)')
plt.ylabel('Sepal Width (cm)')
plt.title(title)
plt.show()

# Plot the decision boundary on the test data
plot_decision_boundary(svm_classifier, X_test, y_test, title='SVM
Decision Boundary')

Accuracy: 1.00
Confusion Matrix:
[[17  0]
 [ 0 13]]

```



SUFYAAN DATASET

```

import pandas as pd
import numpy as np

# Function to generate synthetic diabetes dataset with 100 rows
def generate_synthetic_dataset(num_rows=100):
    np.random.seed(42)

    # Generate random values for each attribute based on some ranges
    age = np.random.randint(20, 70, num_rows)
    pregnancies = np.random.randint(0, 15, num_rows)
    glucose = np.random.randint(60, 200, num_rows)
    blood_pressure = np.random.randint(50, 100, num_rows)
    skin_thickness = np.random.randint(10, 50, num_rows)
    insulin = np.random.randint(50, 200, num_rows)
    bmi = np.random.uniform(20, 40, num_rows)
    diabetes_pedigree = np.random.uniform(0, 1, num_rows)

    # Create the target variable 'Diabetes'
    diabetes = np.random.choice([0, 1], num_rows)

    # Create the synthetic dataset
    dataset = pd.DataFrame({
        'Age': age,
        'Pregnancies': pregnancies,
        'Glucose': glucose,
        'BloodPressure': blood_pressure,
        'SkinThickness': skin_thickness,
        'Insulin': insulin,
        'BMI': bmi,
        'DiabetesPedigreeFunction': diabetes_pedigree,
        'Diabetes': diabetes
    })

    return dataset

# Generate the synthetic dataset with 100 rows
dataset = generate_synthetic_dataset(num_rows=100)

# Save the dataset to a CSV file
dataset.to_csv('synthetic_diabetes_data.csv', index=False)

```